

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

In [2]: df = pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs_per_year\
df
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PI
0	2010-03-01 01:00:00	NaN	0.29	NaN	NaN	NaN	25.090000	29.219999	NaN	68.930000	I
1	2010-03-01 01:00:00	NaN	0.27	NaN	NaN	NaN	24.879999	30.040001	NaN	NaN	I
2	2010-03-01 01:00:00	NaN	0.28	NaN	NaN	NaN	17.410000	20.540001	NaN	72.120003	I
3	2010-03-01 01:00:00	0.38	0.24	1.74	NaN	0.05	15.610000	21.080000	NaN	72.970001	19.410
4	2010-03-01 01:00:00	0.79	NaN	1.32	NaN	NaN	21.430000	26.070000	NaN	NaN	24.670
...	
209443	2010-08-01 00:00:00	NaN	0.55	NaN	NaN	NaN	125.000000	219.899994	NaN	25.379999	I
209444	2010-08-01 00:00:00	NaN	0.27	NaN	NaN	NaN	45.709999	47.410000	NaN	NaN	51.259
209445	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	0.24	46.560001	49.040001	NaN	46.250000	I
209446	2010-08-01 00:00:00	NaN	NaN	NaN	NaN	NaN	46.770000	50.119999	NaN	77.709999	I
209447	2010-08-01 00:00:00	0.92	0.43	0.71	NaN	0.25	76.330002	88.190002	NaN	52.259998	47.150

209448 rows × 17 columns

```
In [3]: df1 = df.fillna(0)
df1
```

Out[3]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	PI
0	2010-03-01 01:00:00	0.00	0.29	0.00	0.0	0.00	25.090000	29.219999	0.0	68.930000	0.000
1	2010-03-01 01:00:00	0.00	0.27	0.00	0.0	0.00	24.879999	30.040001	0.0	0.000000	0.000
2	2010-03-01 01:00:00	0.00	0.28	0.00	0.0	0.00	17.410000	20.540001	0.0	72.120003	0.000
3	2010-03-01 01:00:00	0.38	0.24	1.74	0.0	0.05	15.610000	21.080000	0.0	72.970001	19.410
4	2010-03-01 01:00:00	0.79	0.00	1.32	0.0	0.00	21.430000	26.070000	0.0	0.000000	24.670
...
209443	2010-08-01 00:00:00	0.00	0.55	0.00	0.0	0.00	125.000000	219.899994	0.0	25.379999	0.000
209444	2010-08-01 00:00:00	0.00	0.27	0.00	0.0	0.00	45.709999	47.410000	0.0	0.000000	51.259
209445	2010-08-01 00:00:00	0.00	0.00	0.00	0.0	0.24	46.560001	49.040001	0.0	46.250000	0.000
209446	2010-08-01 00:00:00	0.00	0.00	0.00	0.0	0.00	46.770000	50.119999	0.0	77.709999	0.000
209447	2010-08-01 00:00:00	0.92	0.43	0.71	0.0	0.25	76.330002	88.190002	0.0	52.259998	47.150

209448 rows × 17 columns



```
In [4]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209448 entries, 0 to 209447
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        209448 non-null  object
1   BEN         209448 non-null  float64
2   CO          209448 non-null  float64
3   EBE         209448 non-null  float64
4   MXY         209448 non-null  float64
5   NMHC        209448 non-null  float64
6   NO_2        209448 non-null  float64
7   NOx         209448 non-null  float64
8   OXY         209448 non-null  float64
9   O_3         209448 non-null  float64
10  PM10        209448 non-null  float64
11  PM25        209448 non-null  float64
12  PXY         209448 non-null  float64
13  SO_2        209448 non-null  float64
14  TCH         209448 non-null  float64
15  TOL         209448 non-null  float64
16  station     209448 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 27.2+ MB
```

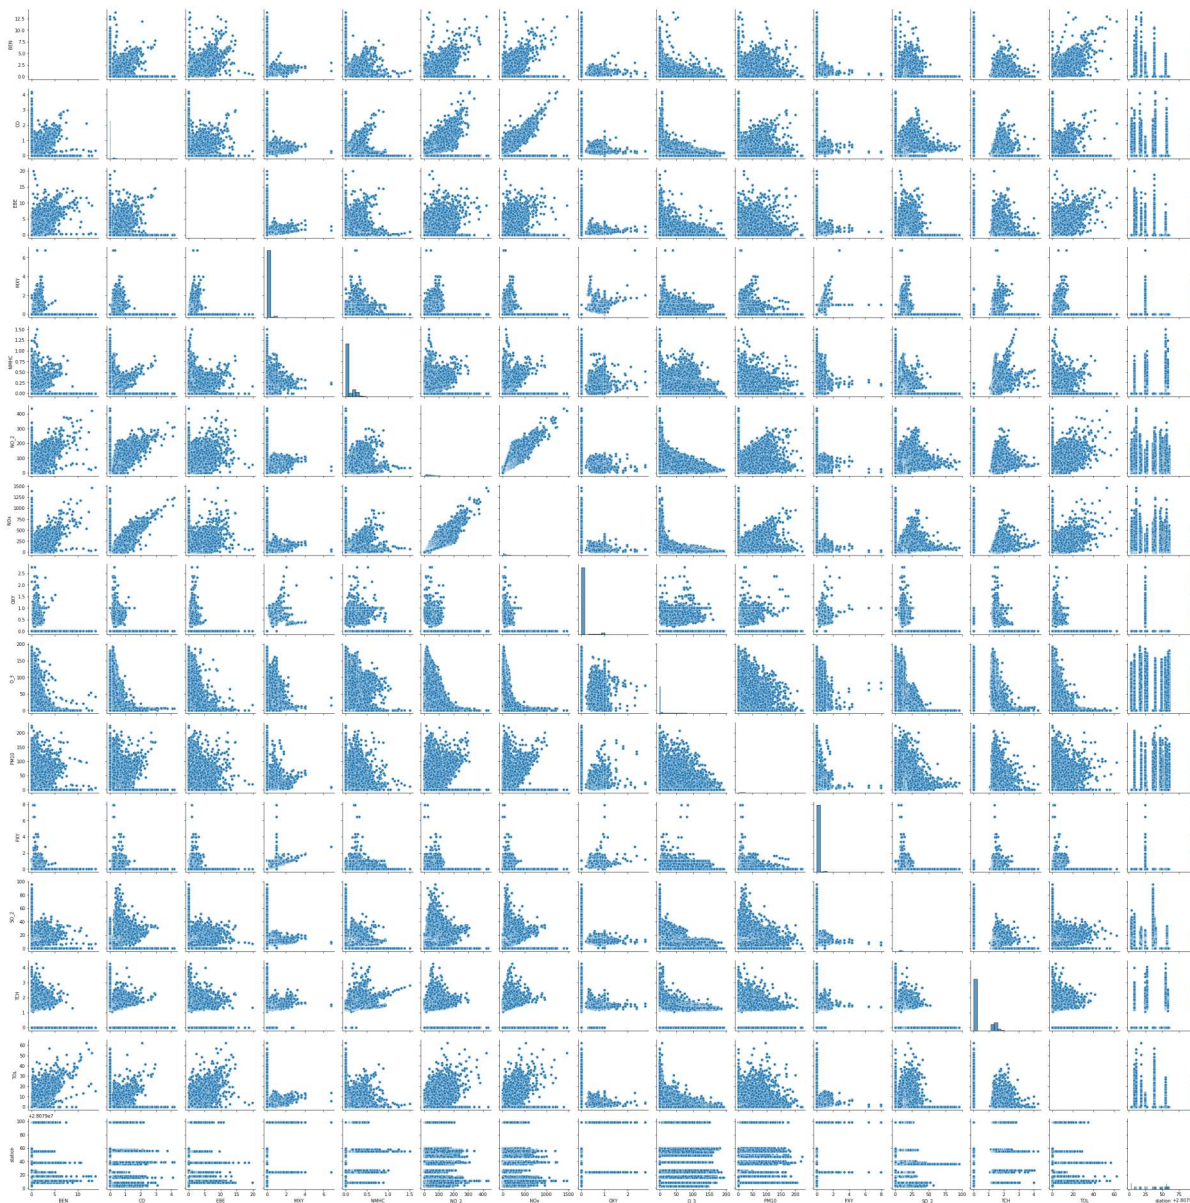
```
In [5]: df1.columns
```

```
Out[5]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
              'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
              dtype='object')
```

```
In [6]: df2 = df1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
                  'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

```
In [7]: sns.pairplot(df2)
```

```
Out[7]: <seaborn.axisgrid.PairGrid at 0x270c83a9100>
```




```
In [40]: sns.distplot(df2['TOL'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

```
Out[40]: <AxesSubplot:xlabel='TOL', ylabel='Density'>
```



```
In [41]: sns.heatmap(df2.corr())
```

```
Out[41]: <AxesSubplot:>
```



Linear Regression

```
In [10]: x = df2[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
               'PM10', 'SO_2', 'PXY', 'TCH']]  
y = df2['TOL']
```

```
In [11]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.30)
```

```
In [12]: from sklearn.linear_model import LinearRegression  
  
lr = LinearRegression()  
lr.fit(x_train,y_train)
```

Out[12]: LinearRegression()

```
In [13]: print(lr.intercept_)  
  
-0.024924509858269328
```


```
In [14]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[14]:

	Co-efficient
BEN	2.209906
CO	-0.516643
EBE	1.208764
MXY	0.409416
NMHC	1.209840
NO_2	0.000958
NOx	0.000925
OXY	-0.712777
O_3	-0.000041
PM10	0.006297
SO_2	0.016920
PXY	0.082471
TCH	-0.389755

```
In [15]: prediction = lr.predict(x_test)
plt.scatter(y_test, prediction)
```

Out[15]: <matplotlib.collections.PathCollection at 0x2708b703bb0>



```
In [16]: print(lr.score(x_test,y_test))
```

0.8014916447995456

```
In [17]: lr.score(x_train,y_train)
```

Out[17]: 0.8055086859693753

Ridge and Lasso

```
In [18]: from sklearn.linear_model import Ridge,Lasso
```

```
In [19]: rr = Ridge(alpha=10)
rr.fit(x_train,y_train)
rr.score(x_train,y_train)
```

Out[19]: 0.8055075396355831

```
In [20]: rr.score(x_test,y_test)
```

Out[20]: 0.8014917172850946

Lasso Regression

```
In [21]: ls = Lasso(alpha=10)
ls.fit(x_train,y_train)
ls.score(x_train,y_train)
```

Out[21]: 0.09156196175090969

```
In [22]: ls.score(x_test,y_test)
```

```
Out[22]: 0.08616910571985181
```

ElasticNET regression

```
In [23]: from sklearn.linear_model import ElasticNet  
es = ElasticNet()  
es.fit(x_train,y_train)
```

```
Out[23]: ElasticNet()
```

```
In [24]: print(es.coef_)
```

```
[ 3.33376457e-01 -0.00000000e+00  6.06608592e-01  0.00000000e+00  
 0.00000000e+00  2.97636043e-03  4.78625298e-03  0.00000000e+00  
 3.46717919e-04  1.97311364e-02  4.16498121e-02  0.00000000e+00  
 0.00000000e+00]
```

```
In [25]: print(es.intercept_)
```

```
-0.20453954044758638
```

```
In [26]: print(es.score(x_test,y_test))
```

```
0.5003170986998571
```

```
In [27]: print(es.score(x_train,y_train))
```

```
0.49927097941131615
```

LogisticRegression

```
In [28]: from sklearn.linear_model import LogisticRegression
```

```
In [29]: feature_matrix = df2.iloc[:,0:15]  
target_vector = df2.iloc[:,-1]
```

```
In [30]: feature_matrix.shape
```

```
Out[30]: (209448, 15)
```

```
In [31]: from sklearn.preprocessing import StandardScaler
```

```
In [32]: fs = StandardScaler().fit_transform(feature_matrix)
```



```
In [33]: logs = LogisticRegression()
logs.fit(fs,target_vector)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:
763: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[33]: LogisticRegression()
```

```
In [34]: observation = [[1.4,1.5,1.6,2.7,2.3,3.3,2.3,4.1,2.3,4.2,1.2,2.1,4.3,6,2.2]]
prediction = logs.predict(observation)
```

```
In [35]: print(prediction)

[28079099]
```

```
In [36]: logs.classes_
```

```
Out[36]: array([28079003, 28079004, 28079008, 28079011, 28079016, 28079017,
                28079018, 28079024, 28079026, 28079027, 28079036, 28079038,
                28079039, 28079040, 28079047, 28079048, 28079049, 28079050,
                28079054, 28079055, 28079056, 28079057, 28079058, 28079059,
                28079060, 28079099], dtype=int64)
```

```
In [37]: from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test = train_test_split(feature_matrix,target_vector,t
```

```
In [38]: print(logs.score(x_test,y_test))

0.04261955916288693
```

```
In [39]: print(logs.score(x_train,y_train))

0.04148336095707748
```

Conclusion

linear regression is bestfit model

linear regression is best fit model for dataset madrid_2001. The score of x_train,y_train is 0.8014916447995456 and x_test and y_test score is 0.8014916447995456.

In []: